

Intrinsically Typed Syntax That Works (Fresh Perspective)

ANONYMOUS AUTHOR(S)

In dependently typed proof assistants, intrinsically typed syntax promises elegant mechanizations: well-typed terms are represented by well-typed data, substitution preserves typing by construction, and type safety proofs collapse to straightforward induction. For simple calculi such as the simply-typed lambda calculus this promise is fulfilled. For realistic polymorphic systems, however, the technique can become impractical: renamings and substitutions indexed by other substitutions force proofs to transport terms across equal types that are not definitionally equal. This phenomenon—informally known as transfer hell—causes even routine metatheoretic arguments to be dominated by administrative reasoning.

This fresh perspective shows how to make intrinsically typed syntax work in practice for polymorphic calculi. We show that giving substitutions canonical computational behavior (via Agda rewriting) yields a mechanization of System F’s metatheory in Agda that avoids explicit transport reasoning.

Our approach simplifies proofs of metatheoretic properties: typing preservation becomes definitional, substitution lemmas and corresponding transfers disappear. Generally, when intrinsic encodings generate transports, the right move is often to strengthen definitional equality rather than add more lemmas. We extract a practical recipe for applying this methodology.

ACM Reference Format:

Anonymous Author(s). 2018. Intrinsically Typed Syntax That Works (Fresh Perspective). *ACM Trans. Program. Lang. Syst.* 37, 4, Article 111 (August 2018), 17 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

Intrinsically typed syntax is widely recommended for mechanizing type systems in dependently typed proof assistants. By representing well-typed terms as well-typed data, typing invariants hold by construction as in a Church-style presentation, so typing preservation is immediate. For small calculi this approach is highly effective. However, for realistic polymorphic calculi such as System F, mechanizations often become complicated: conceptually simple proofs become cluttered with transports due to administrative equalities.

The simply-typed lambda calculus is the canonical example where intrinsically typed syntax is particularly effective. Figure 1 shows typical Agda definitions of the syntax of types, type contexts, variables, and expressions, which are indexed by contexts and types (left column). Building on these definitions we can define a small-step operational semantics and prove type safety. This is because subject reduction holds by construction and progress follows by a simple inductive proof (right column). We elide the definition of the substitution operator $[_]$, as we discuss substitutions in detail in Section 2.1.

The intrinsic encoding makes type safety proofs so concise that they can serve as textbook examples [5]. Nothing comparable to transport reasoning arises in this development as only definitional properties of substitutions are required. However, the same approach breaks down for polymorphic calculi and the remainder of this paper explains how to recover some of its simplicity.

Transfer hell is the colloquial term for being forced to apply type-changing lemmas to terms that are already known to be equal. This phenomenon appears, for example, in intrinsic mechanizations of System F: expression substitutions depend on type substitutions and type substitutions depend on renamings; composing them yields types that are equal only propositionally, forcing transports throughout proofs. The root problem is that substitutions that are propositionally equal are not definitionally equal, so terms must repeatedly be transported across equal types.

```

50 data Type : Set where
51   11      : Type
52   _⇒_ : Type → Type → Type
53
54 data Ctx : Set where
55   0      : Ctx
56   _▷_ : Ctx → Type → Ctx
57
58 data _⊃_ : Ctx → Type → Set where
59   here : (Γ ▷ T) ⊃ T
60   there : Γ ⊃ T → (Γ ▷ U) ⊃ T
61
62 data _⊢_ Γ : Type → Set where
63   con : Γ ⊢ 11
64   var : Γ ⊃ T → Γ ⊢ T
65   lam : (Γ ▷ T) ⊢ U → Γ ⊢ (T ⇒ U)
66   app : Γ ⊢ (T ⇒ U) → Γ ⊢ T → Γ ⊢ U
67
68
69
70 data _→_ {Γ} {T} : Γ ⊢ T → Γ ⊢ T → Set where
71   β-lam : app (lam e1) e2 → (e1 [ e2 ])
72   ξ-app : e1 → e1' → app e1 e2 → app e1' e2
73
74 data Value {Γ} : Γ ⊢ T → Set where
75   con : Value con
76   lam : (e : (Γ ▷ T) ⊢ U) → Value (lam e)
77
78 data Progress (e : Γ ⊢ T) : Set where
79   done : Value e → Progress e
80   step : e → e' → Progress e
81
82 progress : (e : 0 ⊢ T) → Progress e
83 progress con      = done con
84 progress (lam e) = done (lam e)
85 progress (app e1 e2)
86   with progress e1
87   ... | step x      = step (ξ-app x)
88   ... | done (lam e) = step β-lam

```

Fig. 1. Call-by-name simply-typed lambda calculus (adapted from the PLFA textbook [5])

One workaround is to state definitions and proofs with explicit equivalence relations, but this practice is cumbersome, inefficient, and far away from the concision of definitional equality. Our proposal thus shifts the focus to computation on substitutions: if extensionally equal substitutions do not normalize identically, transport reasoning is unavoidable.

Hence, the simple message of this fresh perspective is: when intrinsically typed syntax becomes complicated, the problem is often not the encoding of terms but the definition of equality on substitutions. Strengthening definitional equality can simplify metatheoretical proofs more effectively than adding automation or additional lemmas.

We substantiate our claims with a case study on renaming and substitution in an intrinsic encoding of System F. We strengthen definitional equality by internalizing the equational theory of substitutions (the σ -calculus) using Agda’s rewriting mechanism. As a consequence, extensionally equal substitutions normalize to the same form and transports disappear from routine proofs.

Contributions

- We show how to eliminate “transfer hell” in intrinsically typed encodings of System F by internalizing the equational laws of the σ -calculus for substitutions using Agda’s rewriting mechanism.
- We demonstrate that the resulting rewrite system is compatible with Agda’s definitional equality (in particular, terminating and confluent) for a functional representation of substitutions.
- We mechanize parts of the metatheory of System F, obtaining substantially simpler proofs and an intrinsic treatment of type conversion for the latter.
- From this experience we extract practical guidelines for designing intrinsically typed syntax in dependently typed proof assistants.

```

99 data Type (n : Nat) : Set where
100   ' _ : Var n → Type n
101   ∀α : Type (1 + n) → Type n
102   _⇒_ : Type n → Type n → Type n
103
104   _ : Type 0 - a closed type:  ∀α. α→α
105   _ = ∀α (' zero ⇒ ' zero)
106   _ : Type 0 - ∀αβ. α→β→α
107   _ = ∀α (∀α (' suc zero ⇒ ' zero ⇒ ' suc zero))

```

Fig. 2. Syntax of System F types (Agda definition left, examples right)

All code fragments in this document are extracted from typechecked Agda code, which is available as a supplement.

Related Work and Positioning

To our knowledge, previous intrinsically typed mechanizations of System F [1, 2, 10] follow a natural design: expression syntax is indexed by typing derivations and substitutions are indexed by substitutions at the type level. While elegant at the level of definitions, this design makes proofs brittle, because they are prone to transfer hell. Our work eliminates transfers resulting from substitutions, thus avoiding transfer hell.

Recent work studies rewrite rules in dependent type theories, including Agda [3] and Rocq [6, 7]. In Rewriting Type Theory (RTT), the reduction relation is extended by turning propositional equalities into reduction rules: a term matched by the left-hand side of an equation reduces to the right-hand side. Rewrites are safe if they preserve subject reduction and consistency. The former is guaranteed by confluence of the combined rewrite system, the latter by only rewriting proved propositional equalities. Our work is based on Agda's implementation of RTT.

Our approach is inspired by the work on Autosubst [11–13]. While they rely on tactics to normalize substitutions, we do so by rewriting [4]. We build on a standard functional representation of renamings and substitutions [1]. The main challenge in setting up the σ -calculus is to fine-tune the reduction behavior of the definitions with the rewrite rules to ensure that the combined reduction relation remains terminating and confluent.

Overview

Section 2 presents our scoped encoding of System F types. It discusses the standard encoding of functional renamings and substitutions (Section 2.2) and introduces the σ -calculus laws (Section 2.3). Section 3 presents the intrinsically typed encoding of System F expressions. It defines expression substitutions (Section 3.2) and a small-step operational semantics (Section 3.3). Section 4 demonstrates a concrete instance of transfer hell, once treated manually and once by using our method. Section 5 outlines methodological issues along with recommendations and guidelines.

2 System F Types

2.1 Syntax

The syntax of System F types uses the standard de Bruijn encoding of variables by numbers. The type `Type n` (Figure 2) stands for the set of types with n type variables in scope. Henceforth, we call such n a *scope*. Type variables for scope n are drawn from the type `Var n`, i.e., the set $\{0, \dots, n-1\}$, written as `zero`, `suc zero`, `suc (suc zero)`, ... and ranged over by α . The $\forall\alpha$ -binder introduces a new type variable, and $\Rightarrow$ is the infix function type constructor. We call this style of syntax *intrinsically scoped* as it can only express well-scoped types. We let $\mathbf{T}$ range over types.

148	- renamings	- substitutions
149	$\text{Ren} : \text{Nat} \rightarrow \text{Nat} \rightarrow \text{Set}$	$\text{Sub} : \text{Nat} \rightarrow \text{Nat} \rightarrow \text{Set}$
150	$\text{Ren } n_1 n_2 = \text{Var } n_1 \rightarrow \text{Var } n_2$	$\text{Sub } n_1 n_2 = \text{Var } n_1 \rightarrow \text{Type } n_2$
151		
152	opaque	opaque
153	- weakening	- lift renaming to substitution
154	$\text{wk}^R : \text{Ren } n (1 + n)$	$\langle _ \rangle : \text{Ren } n_1 n_2 \rightarrow \text{Sub } n_1 n_2$
155	$\text{wk}^R = \text{suc}$	$\langle \zeta \rangle \alpha = \langle \alpha \ \&^R \ \zeta \rangle$
156		
157	- identity renaming	- extend with new type
158	$\text{id}^R : \text{Ren } n n$	$_ \cdot^S _ : \text{Type } n_2 \rightarrow \text{Sub } n_1 n_2 \rightarrow \text{Sub } (1 + n_1) n_2$
159	$\text{id}^R \alpha = \alpha$	$(T \cdot^S \eta) \text{zero} = T$
160		$(T \cdot^S \eta) (\text{suc } \alpha) = \eta \alpha$
161	- extend with new variable	- apply substitution to variable
162	$_ \cdot^R _ : \text{Var } n_2 \rightarrow \text{Ren } n_1 n_2 \rightarrow \text{Ren } (1 + n_1) n_2$	$_ \&^S _ : \text{Var } n_1 \rightarrow \text{Sub } n_1 n_2 \rightarrow \text{Type } n_2$
163	$(\alpha \cdot^R \zeta) \text{zero} = \alpha$	$\alpha \ \&^S \eta = \eta \alpha$
164	$(_ \cdot^R \zeta) (\text{suc } \alpha) = \zeta \alpha$	
165	- apply renaming to variable	- lifting
166	$_ \&^R _ : \text{Var } n_1 \rightarrow \text{Ren } n_1 n_2 \rightarrow \text{Var } n_2$	$_ \uparrow^S : \text{Sub } n_1 n_2 \rightarrow \text{Sub } (1 + n_1) (1 + n_2)$
167	$\alpha \ \&^R \zeta = \zeta \alpha$	$_ \uparrow^S \eta = (\langle \text{zero} \rangle \cdot^S \lambda \alpha \rightarrow (\eta \alpha) [\text{wk}^R]^R)$
168		
169	- left-to-right composition	- apply substitution to type
170	$_ \circ^R _ : \text{Ren } n_1 n_2 \rightarrow \text{Ren } n_2 n_3 \rightarrow \text{Ren } n_1 n_3$	$_ _]^S : \text{Type } n_1 \rightarrow \text{Sub } n_1 n_2 \rightarrow \text{Type } n_2$
171	$(\zeta_1 \circ^R \zeta_2) \alpha = \zeta_2 (\zeta_1 \alpha)$	$(\langle \alpha \rangle [\eta]^S = \alpha \ \&^S \eta$
172		$(\forall \alpha T) [\eta]^S = \forall \alpha (T [\eta \uparrow^S]^S)$
173	- lifting	$(T_1 \Rightarrow T_2) [\eta]^S = (T_1 [\eta]^S) \Rightarrow (T_2 [\eta]^S)$
174	$_ \uparrow^R : \text{Ren } n_1 n_2 \rightarrow \text{Ren } (1 + n_1) (1 + n_2)$	
175	$_ \uparrow^R \zeta = \text{zero} \cdot^R (\zeta \circ^R \text{wk}^R)$	opaque
176		- left-to-right composition
177	- apply renaming to type	$_ \circ^S _ : \text{Sub } n_1 n_2 \rightarrow \text{Sub } n_2 n_3 \rightarrow \text{Sub } n_1 n_3$
178	$_ _]^R : \text{Type } n_1 \rightarrow \text{Ren } n_1 n_2 \rightarrow \text{Type } n_2$	$(\eta_1 \circ^S \eta_2) \alpha = (\eta_1 \alpha) [\eta_2]^S$
179	$(\langle \alpha \rangle [\zeta]^R = \langle \alpha \ \&^R \ \zeta \rangle$	
180	$(\forall \alpha T) [\zeta]^R = \forall \alpha (T [\zeta \uparrow^R]^R)$	
181	$(T_1 \Rightarrow T_2) [\zeta]^R = (T_1 [\zeta]^R) \Rightarrow (T_2 [\zeta]^R)$	

Fig. 3. Renamings and substitutions on types

2.2 Renaming and Substitution

The most important operations on types are consistent renaming of type variables and substitution. Figure 3 contains the first step in defining renamings (left column) and substitutions (right column). Our definitions of renamings and substitutions in functional style are standard [1, 5, 9]. For now, we ignore the use of `opaque`. We defer the discussion of using it to Section 2.3.

Table 1 summarizes the operator layer used throughout the development. The table is intended as a reference map: renaming operators $\zeta : \text{Ren } n_1 n_2$ capture structure-preserving variable reindexing from scope n_1 to scope n_2 , while substitution operators $\eta : \text{Sub } n_1 n_2$ perform capture-avoiding substitution of types from scope n_2 for variables in a type at scope n_1 . Rather than introducing each symbol ad hoc in later proofs, we fix this interface once and reuse it uniformly.

Family	Operator	Role
Renaming	wk^R	weakening renaming
Renaming	id^R	identity renaming
Renaming	$\alpha \cdot^R \zeta$	extend renaming with head variable α
Renaming	$\alpha \&^R \zeta$	apply renaming to a variable
Renaming	$\zeta_1 \circ^R \zeta_2$	compose renamings
Renaming	$\zeta \uparrow^R$	lift renaming under a binder
Renaming	$T[\zeta]^R$	apply renaming to a type
Substitution	$\langle \zeta \rangle$	embed renaming as substitution
Substitution	$T \cdot^S \eta$	extend substitution with head type T
Substitution	$\alpha \&^S \eta$	apply substitution to a variable
Substitution	$\eta_1 \circ^S \eta_2$	compose substitutions
Substitution	$\eta \uparrow^S$	lift substitution under a binder
Substitution	$T[\eta]^S$	apply substitution to a type

Table 1. Operator inventory for type renamings and substitutions.

Two design choices are important for what follows. First, renamings and substitutions are represented functionally, which keeps composition and lifting explicit. Second, the operators in this interface are the symbols on which we later impose the σ -calculus equations as rewrite rules. This separation between interface and equations enables us to control reduction behavior precisely.

2.3 The σ -Calculus with First-class Renamings

This section presents the equalities of the σ -calculus. These are the equalities that we *actually* use for computing with renamings and substitutions, rather than their defining equations in Figure 3. The **opaque** blocks control this behavior. Function symbols defined in an **opaque** block are usable outside the block, but they do not reduce. The only reducing definitions in the renaming block are $\zeta \uparrow^R$ for lifting a renaming and application of a renaming to a type. In the substitution block, only the application of a substitution to a type reduces.

This selective reduction behavior is essential for our approach because we want to compute using the laws of the σ -calculus with first-class renamings on top of the thus defined symbols, rather than their definitions. Blocking their reduction is necessary because the equalities from σ -calculus interfere adversely with the standard definitional equalities. The original defining equalities are available for proofs: We can open the opaque definitions locally and proceed as usual.

Next we introduce equational laws of the σ -calculus with first-class renamings [11]. All of these equalities are supported by proofs using the definitions in Figure 3. The proofs are available in the supplement, but they are not included here for space reasons.

The first group presents computation rules for substitutions.

$$\begin{aligned}
\text{beta-ext-zero} &: \text{zero} \&^S (T \cdot^S \eta) \equiv T \\
\text{beta-ext-suc} &: \text{suc } \alpha \&^S (T \cdot^S \eta) \equiv \alpha \&^S \eta \\
\text{beta-rename} &: \alpha \&^S \langle \zeta \rangle \equiv (\alpha \&^R \zeta) \\
\text{beta-comp} &: \alpha \&^S (\eta_1 \circ^S \eta_2) \equiv (\alpha \&^S \eta_1) [\eta_2]^S \\
\text{beta-lift} &: \eta \uparrow^S \equiv (\text{zero}) \cdot^S (\eta \circ^S \langle \text{wk}^R \rangle)
\end{aligned}$$

Starting with the application of a substitution to a variable, there is one equation for each possible form of a substitution: extended substitution, lifted renaming, and composition of substitutions. The final equation in this group characterizes lifting in terms of extension, weakening, and composition. It is *not* possible to use this equation to define lifting in Figure 3 as it would lead to circularity.

The second group contains laws about interactions between substitutions.

$$\begin{aligned}
&\text{associativity} : (\eta_1 \circ^S \eta_2) \circ^S \eta_3 && \equiv \eta_1 \circ^S (\eta_2 \circ^S \eta_3) \\
&\text{distributivity} : (T \cdot^S \eta_1) \circ^S \eta_2 && \equiv (T [\eta_2]^S) \cdot^S (\eta_1 \circ^S \eta_2) \\
&\text{interact} : \langle \text{wk}^R \rangle \circ^S (T \cdot^S \eta) && \equiv \eta \\
&\text{comp-id}_r : \eta \circ^S \langle \text{id}^R \rangle && \equiv \eta \\
&\text{comp-id}_l : \langle \text{id}^R \rangle \circ^S \eta && \equiv \eta \\
&\eta\text{-id} : (\text{zero } \{n\}) \cdot^S \langle \text{wk}^R \rangle && \equiv \langle \text{id}^R \rangle \\
&\eta\text{-law} : (\text{zero } \&^S \eta) \cdot^S (\langle \text{wk}^R \rangle \circ^S \eta) && \equiv \eta
\end{aligned}$$

Composition is associative and it distributes over extension. First weakening and then extending with T has no effect, as T is substituted for a non-existing variable. The identity substitution is a right and left identity. We conclude this group with two η -laws about the interaction of extension and weakening.¹

There are analogous rules for renamings corresponding to the first two groups. These rules are omitted for space reasons.

The third group presents laws that govern the composition of different kinds of term traversal, that is, the behavior of the identity renaming and the four possible compositions of renamings and substitutions.

$$\begin{aligned}
&\text{identity}_r : T [\text{id}^R]^R && \equiv T \\
&\text{compositionality}^{RS} : (T [\zeta_1]^R) [\eta_2]^S && \equiv T [\langle \zeta_1 \rangle \circ^S \eta_2]^S \\
&\text{compositionality}^{RR} : (T [\zeta_1]^R) [\zeta_2]^R && \equiv T [\zeta_1 \circ^R \zeta_2]^R \\
&\text{compositionality}^{SR} : (T [\eta_1]^S) [\zeta_2]^R && \equiv T [\eta_1 \circ^S \langle \zeta_2 \rangle]^S \\
&\text{compositionality}^{SS} : (T [\eta_1]^S) [\eta_2]^S && \equiv T [\eta_1 \circ^S \eta_2]^S
\end{aligned}$$

The final group contains laws that mediate between substitutions and renamings. Substitution of a lifted renaming is a renaming; the composition of two lifted renamings is the lifting of their composition as renamings.

$$\begin{aligned}
&\text{coincidence} : T [\langle \zeta \rangle]^S && \equiv T [\zeta]^R \\
&\text{coincidence-comp} : \langle \zeta_1 \rangle \circ^S \langle \zeta_2 \rangle && \equiv \langle \zeta_1 \circ^R \zeta_2 \rangle
\end{aligned}$$

In general, the extraction of renamings from substitutions requires a dedicated solving strategy [13] and we have omitted some further coincidence laws for brevity.

Orienting these equations from left to right results in a confluent rewrite system for the σ -calculus that we install in Agda's computation engine using the **REWRITE** pragma [4]. Confluence of these rewrite equations is mechanically checked by Agda using the flags `-rewriting` `-confluence-check`. We discuss safety considerations for rewriting in Section 5.1. From now on, all equations shown in this section Section 2.3 are treated as definitional equalities.

As a consistency check for the definitions, we prove in Figure 4 that lifting a substitution and the action of a substitution satisfy the functor laws. With σ -calculus rewriting installed, all these lemmas hold definitionally as evidenced by their proof using reflexivity (**refl**). The lemmas marked with $*$ correspond directly to laws from the σ -calculus.² The lifting laws are shown on the left, the application laws on the right. The analogous results for renamings follow exactly the same pattern.

¹In the statement of η -id, we write `zero {n}` to tell Agda's type checker that n is the scope of type variable `zero`. This notation supplies an *implicit argument* that Agda can usually infer.

²The naming of the lemmas is taken from Chapman et al. [2] to enable direct comparison.

```

295 lifts*-id : (idS {n} ↑S) ≡ idS
296 lifts*-id = refl
297
298 lifts*-comp : (η' ∘S η) ↑S ≡ (η' ↑S) ∘S (η ↑S)
299 lifts*-comp = refl
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343

```

$$\text{sub}^*\text{-id} : T[\text{id}^S] \equiv T$$

$$\text{sub}^*\text{-id} = \text{refl}$$

$$\text{sub}^*\text{-var} : (\alpha) [\eta]^S \equiv \alpha \&^S \eta$$

$$\text{sub}^*\text{-var} = \text{refl} - *$$

$$\text{sub}^*\text{-comp} : T[\eta \circ^S \eta']^S \equiv (T[\eta]^S) [\eta']^S$$

$$\text{sub}^*\text{-comp} = \text{refl} - *$$

Fig. 4. Functor laws for substitution proved

```

306
307 weaken : Type n → Type (1 + n)
308 weaken T = T [ wkR ]R
309
310 _[_]* : Type (1 + n) → Type n → Type n
311 T [ T' ]* = T [ T' .S idS ]S
312
313 data Ctx : Nat → Set where
314   () : Ctx zero
315   _▷_ : Ctx n → Type n → Ctx n
316   _▷* : Ctx n → Ctx (1 + n)
317
318 data _≡_ : Ctx n → Type n → Set where
319   zero : (Γ ▷ T) ≡ T
320   suc : Γ ≡ T → (Γ ▷ T') ≡ T'
321   suc* : Γ ≡ T → (Γ ▷*) ≡ weaken T
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343

```

$$\text{data Expr } (\Gamma : \text{Ctx } n) : \text{Type } n \rightarrow \text{Set where}$$

$$_ : \Gamma \ni T \rightarrow \text{Expr } \Gamma T$$

$$\lambda x : \text{Expr } (\Gamma \triangleright T_1) T_2 \rightarrow \text{Expr } \Gamma (T_1 \Rightarrow T_2)$$

$$_ \cdot _ : \text{Expr } \Gamma (T_1 \Rightarrow T_2) \rightarrow \text{Expr } \Gamma T_1 \rightarrow \text{Expr } \Gamma T_2$$

$$\Lambda \alpha : \text{Expr } (\Gamma \triangleright^*) T \rightarrow \text{Expr } \Gamma (\forall \alpha T)$$

$$_ \cdot^* _ : \text{Expr } \Gamma (\forall \alpha T) \rightarrow (T' : \text{Type } n) \rightarrow \text{Expr } \Gamma (T [T']^*)$$

Fig. 5. Type contexts, variables, and expressions

3 System F Expressions

3.1 Syntax

Figure 5 defines the ingredients of intrinsically typed expressions: type weakening (**weaken**) and single substitution $T [T']^*$ of a type T' for the innermost variable of T along with type contexts Γ , expression variables x , and expression syntax e . These ingredients provide the infrastructure for defining substitution and semantics in the following subsections.

Type contexts and expression variables follow the work of Chapman et al. [2]. A type context $\Gamma : \text{Ctx } n$ tracks expression and type variables in scope n . The operator $_ \triangleright _$ adds a value binding while $_ \triangleright^*$ adds a type binding to Γ . Correspondingly, expression variables are implemented as typed de Bruijn indices with an extra case **suc*** to skip over type variables in the context.

Skipping with **suc*** requires weakening of the type as it is used under one more type binding. The lambda expression adds a value binding, while the type lambda adds a type binding. Type application requires the single substitution.

Expressions $e : \text{Expr } \Gamma T$ are indexed by a context Γ and a type T with the intended reading as a typing judgment $\Gamma \vdash e : T$. As customary in an intrinsically typed setting, each expression constructor corresponds to a typing rule in System F. Figure 6 presents some example expressions

```

344  $\mathbb{N}^c : \text{Type } 0 *$ 
345  $\mathbb{N}^c = \forall \alpha ((\alpha \Rightarrow \alpha) \Rightarrow (\alpha \Rightarrow \alpha))$ 
346
347  $\text{one}^c : \text{Expr } 0 \mathbb{N}^c$ 
348  $\text{one}^c = \Lambda \alpha (\lambda g (\lambda x ('g \cdot 'x)))$ 
349
350  $\text{succ}^c : \text{Expr } 0 (\mathbb{N}^c \Rightarrow \mathbb{N}^c)$ 
351  $\text{succ}^c = \lambda x (\Lambda \alpha (\lambda g (\lambda x ('g \cdot (((('(\text{succ}(\text{succ}(\text{succ}^* \text{zero})))) \cdot * ' \alpha) \cdot 'g) \cdot 'x))))))$ 
352
353
354
355

```

Fig. 6. Church numerals in System F

and types. For readability, we have aliases α for $'Z$, β for $'(S Z)$, and analogously for expression variables and binders.

3.2 Renaming and Substitution

The definition of a small-step operational semantics requires expression renamings and substitution to implement beta reduction for value abstractions and type abstractions. Renamings and substitutions are both indexed by type renamings and type substitutions, respectively.

Expression renamings ρ are indexed by type renamings ζ . As usual, they map a variable from one context to another. Notationally, we write $\zeta \mid \dots$ for any operation on renamings that is indexed by type renaming ζ . We only show the following representative definition as an example. The remaining operations may be found in the supplement.

```

367  $\_ \mid \_ \Rightarrow^R \_ : \text{Ren } n_1 n_2 \rightarrow \text{Ctx } n_1 \rightarrow \text{Ctx } n_2 \rightarrow \text{Set}$ 
368  $\zeta \mid \Gamma_1 \Rightarrow^R \Gamma_2 = \forall T \rightarrow (x : \Gamma_1 \ni T) \rightarrow \Gamma_2 \ni (T \mid \zeta ]^R)$ 
369

```

Expression substitutions σ are indexed by type substitutions η . A substitution takes a variable x of type T in context Γ_1 and maps it to an expression typed in context Γ_2 adjusting its type according to the type substitution: $T \mid \eta]^S$. The notation is analogous to the one for renamings.

```

373  $\_ \mid \_ \Rightarrow^S \_ : \text{Sub } n_1 n_2 \rightarrow \text{Ctx } n_1 \rightarrow \text{Ctx } n_2 \rightarrow \text{Set}$ 
374  $\eta \mid \Gamma_1 \Rightarrow^S \Gamma_2 = \forall T \rightarrow (x : \Gamma_1 \ni T) \rightarrow \text{Expr } \Gamma_2 (T \mid \eta ]^S)$ 
375

```

Applying a substitution to a variable corresponds to function application. Nevertheless, we use the same interface as for type substitutions to enable using the same rewriting machinery on the expression level again. For this definition, we include the typing of all symbols, which is similar for all further definitions and omitted because it can be inferred.

```

380  $\_ \mid \_ \&^S \_ : \forall \{\Gamma_1 : \text{Ctx } n_1\} \{\Gamma_2 : \text{Ctx } n_2\} (\eta : \text{Sub } n_1 n_2)$ 
381  $\rightarrow (x : \Gamma_1 \ni T) \rightarrow (\sigma : \eta \mid \Gamma_1 \Rightarrow^S \Gamma_2) \rightarrow \text{Expr } \Gamma_2 (T \mid \eta ]^S)$ 
382  $\eta \mid x \&^S \sigma = \sigma \_ x$ 
383

```

The identity substitution for expressions appears simple only because of the installed rewriting of type substitutions. Without the rewriting, the right-hand side's type is T , but the expected type is $T \mid \text{id}^S]^S$, which would require the transport lemma `comp-idr`.

```

387  $\text{id}^S : \text{id}^S \mid \Gamma \Rightarrow^S \Gamma$ 
388  $\text{id}^S \_ = \lambda x \rightarrow 'x$ 
389

```

Extending a substitution is standard: given an expression e and a substitution σ , we return a substitution that sends variable `zero` to e and otherwise resorts to σ .

$_ _ \cdot^S _ : \forall \eta \rightarrow (e : \text{Expr } \Gamma_2 (T[\eta]^S)) \rightarrow (\sigma : \eta \mid \Gamma_1 \Rightarrow^S \Gamma_2) \rightarrow \eta \mid (\Gamma_1 \triangleright T) \Rightarrow^S \Gamma_2$
 $(_ \mid e \cdot^S \sigma) _ \text{zero} = e$
 $(_ \mid e \cdot^S \sigma) _ (\text{suc } x) = \sigma _ x$

Due to the structure of contexts, we need two lifting lemmas when defining the application of a substitution to an expression, one for value bindings and another for type bindings. In both cases, the substitution σ needs to be adjusted (i.e., lifted) to accommodate the newly introduced binding. The former is standard: it maps the newly introduced variable to itself (**zero**) and weakens the images of σ accordingly. However, the definition of lifting requires a transport lemma to line up the type substitutions, which is applied by rewriting. The last part of the definition $\text{id}^R \mid \dots \text{Wk}^R \dots$ applies an expression renaming (weakening by $T[\eta]^S$) indexed by id^R to the output of σ .

$_ _ \uparrow^S _ : \forall \eta \rightarrow (\sigma : \eta \mid \Gamma_1 \Rightarrow^S \Gamma_2) \rightarrow \forall T \rightarrow \eta \mid (\Gamma_1 \triangleright T) \Rightarrow^S (\Gamma_2 \triangleright (T[\eta]^S))$
 $\eta \mid \sigma \uparrow^S T = \eta \mid (\text{' zero}) \cdot^S \lambda _ x \rightarrow \text{id}^R \mid (\sigma _ x) [\text{Wk}^R (T[\eta]^S)]^R$

Lifting a substitution over a type binding is non-standard. The type substitution requires lifting, but no rewriting is needed.

$_ _ \uparrow^{S*} : \forall \eta \rightarrow (\sigma : \eta \mid \Gamma_1 \Rightarrow^S \Gamma_2) \rightarrow (\eta \uparrow^S) \mid (\Gamma_1 \triangleright^*) \Rightarrow^S (\Gamma_2 \triangleright^*)$
 $(\eta \mid \sigma \uparrow^{S*}) = (\eta \uparrow^S \langle \text{wk}^R \rangle) \mid (\text{' zero}) \cdot^{S*} \lambda _ x \rightarrow \text{wk}^R \mid (\sigma _ x) [\text{wk}^{R*}]^R$

Substitution application is defined by structural recursion on expressions. The cases dispatch to the operations that we introduced: applying a substitution to a variable, lifting over a value binding (case λx), lifting over a type binding (case $\Lambda \alpha$), the cases for application and type-application are homomorphic.

$_ _ _ _ _ : (\eta : \text{Sub } n_1 \ n_2) \rightarrow (e : \text{Expr } \Gamma_1 \ T) \rightarrow (\sigma : \eta \mid \Gamma_1 \Rightarrow^S \Gamma_2) \rightarrow \text{Expr } \Gamma_2 (T[\eta]^S)$
 $\eta \mid (\text{' } x) [\sigma]^S = \eta \mid x \&^S \sigma$
 $\eta \mid (\lambda x \ e) [\sigma]^S = \lambda x \ (\eta \mid e [\eta \mid \sigma \uparrow^S _]^S)$
 $\eta \mid (\Lambda \alpha \ e) [\sigma]^S = \Lambda \alpha \ ((\eta \uparrow^S) \mid e [\eta \mid \sigma \uparrow^{S*}]^S)$
 $\eta \mid (e \cdot e_1) [\sigma]^S = (\eta \mid e [\sigma]^S) \cdot (\eta \mid e_1 [\sigma]^S)$
 $\eta \mid (e \cdot^* T) [\sigma]^S = (\eta \mid e [\sigma]^S) \cdot^* (T[\eta]^S)$

3.3 Semantics

We are now ready to define a small-step operational semantics. As in the introduction, we opt for call-by-name for concision. The definition of single substitution, needed for beta reduction, is analogous to the one on types: we build the substitution by extending the identity substitution with the argument e' . As usual in an intrinsically typed setting, this definition already proves that substitution preserves typing.

$_ _ _ : \text{Expr } (\Gamma \triangleright T') \ T \rightarrow \text{Expr } \Gamma \ T' \rightarrow \text{Expr } \Gamma \ T$
 $e [\ e'] = \text{id}^S \mid e [\text{id}^S \mid e' \cdot^S \text{id}^S]^S$

In addition, we have to substitute a type into an expression to cater for beta reduction of type applications. This substitution is also implemented by an expression substitution indexed by a type substitution that performs the actual change.

$_ _ _ _ _ : \text{Expr } (\Gamma \triangleright^*) \ T \rightarrow (T' : \text{Type } n) \rightarrow \text{Expr } \Gamma \ (T[\ T']^*)$
 $e [\ T']^* = (T' \cdot^S \text{id}^S) \mid e [\text{id}^S \mid T' \cdot^{S*} \text{id}^S]^S$

Due to the extended type substitution on the left, we also have to adapt the expression substitution on the right to the extended context structure. This adaptation is the type-level analogue to extension:

```

442 data Value : Expr  $\Gamma$   $T \rightarrow$  Set where
443    $\lambda x : (e : \text{Expr } (\Gamma \triangleright T_1) T_2) \rightarrow \text{Value } (\lambda x e)$ 
444    $\Lambda \alpha : (v : \text{Value } e) \rightarrow \text{Value } (\Lambda \alpha e)$ 
445
446 data Progress : Expr  $\Gamma$   $T \rightarrow$  Set where
447   done : (v : Value e)  $\rightarrow$  Progress e
448   step : (e  $\rightarrow$  e' : e  $\rightarrow$  e')  $\rightarrow$  Progress e
449
450 NoVar : Ctx n  $\rightarrow$  Set
451 NoVar  $\Gamma = \forall \{T\} \rightarrow \neg (\Gamma \ni T')$ 
452
453 progress : NoVar  $\Gamma \rightarrow (e : \text{Expr } \Gamma T) \rightarrow \text{Progress } e$ 
454 progress nv (' x) =  $\perp$ -elim (nv x)
455 progress nv ( $\lambda x e$ ) = done ( $\lambda x e$ )
456 progress nv (e  $\cdot$  e1)
457   with progress nv e
458   ... | done ( $\lambda x e_2$ ) = step  $\beta$ - $\lambda$ 
459   ... | step e  $\rightarrow$  e' = step ( $\xi$ - $\cdot$  e  $\rightarrow$  e')
460 progress nv ( $\Lambda \alpha e$ )
461   with progress ( $\lambda \{ (\text{suc}^* x) \rightarrow \text{nv } x \}$ ) e
462   ... | done v = done ( $\Lambda \alpha v$ )
463   ... | step e  $\rightarrow$  e' = step ( $\xi$ - $\Lambda$  e  $\rightarrow$  e')
464 progress nv (e  $\cdot$  T')
465   with progress nv e
466   ... | done ( $\Lambda \alpha v$ ) = step  $\beta$ - $\Lambda$ 
467   ... | step e  $\rightarrow$  e' = step ( $\xi$ - $\cdot$  e  $\rightarrow$  e')

```

Fig. 7. Progress for System F

```

462  $\_ \cdot \_ . S^* \_ : \forall \eta T \rightarrow (\sigma : \eta \mid \Gamma_1 \Rightarrow^S \Gamma_2) \rightarrow (T \cdot^S \eta) \mid (\Gamma_1 \triangleright^* ) \Rightarrow^S \Gamma_2$ 
463 ( $\_ \mid T \cdot S^* \sigma$ )  $\_$  (suc* x) =  $\sigma \_ x$ 

```

With substitution in place, we present the definition of the call-by-name small-step reduction relation. As usual, such a definition of reduction for intrinsically-typed syntax implies subject reduction by construction.

```

468 data  $\_ \rightarrow \_ : \text{Expr } \Gamma T \rightarrow \text{Expr } \Gamma T \rightarrow \text{Set}$  where
469    $\beta$ - $\lambda : (\lambda x e_1 \cdot e_2) \rightarrow (e_1 [e_2])$ 
470    $\beta$ - $\Lambda : (\Lambda \alpha e \cdot^* T') \rightarrow (e [^* T' ^*])$ 
471    $\xi$ - $\cdot : e_1 \rightarrow e_1' \rightarrow (e_1 \cdot e_2) \rightarrow (e_1' \cdot e_2)$ 
472    $\xi$ - $\cdot^* : e \rightarrow e' \rightarrow (e \cdot^* T) \rightarrow (e' \cdot^* T)$ 
473    $\xi$ - $\Lambda : e \rightarrow e' \rightarrow (\Lambda \alpha e) \rightarrow (\Lambda \alpha e')$ 

```

To obtain type safety, it remains to prove progress. Figure 7 presents the complete proof of progress for call-by-name System F using the structure from PLFA [5]. A value is either a lambda or a type lambda where the body is already a value. An expression satisfies progress if it is either a value (done) or if it can make a step (step).

Due to the definition of value for type lambdas, we have to establish progress for contexts that may contain type bindings, but no value bindings. These contexts are characterized literally by the function NoVar.

The proof of progress is by induction on the structure of expressions. The overall structure of this proof closely follows Chapman et al [2]. Their proof is more complicated because their semantics is call-by-value, their language supports isorecursive types, and their subject language is System F ω .

Next, we move on to study an example of transfer hell in detail.

4 Transfer Heaven and Transfer Hell

In this section, we study functorial and monadic properties of expression substitutions. These properties are needed once we consider properties of logical relations [10].

$$\begin{aligned} _ , _ \stackrel{S}{\circlearrowleft} _ : & \forall \eta_1 \eta_2 \rightarrow (\sigma_1 : \eta_1 \mid \Gamma_1 \Rightarrow^S \Gamma_2) \rightarrow (\sigma_2 : \eta_2 \mid \Gamma_2 \Rightarrow^S \Gamma_3) \rightarrow (\eta_1 \stackrel{S}{\circlearrowleft} \eta_2) \mid \Gamma_1 \Rightarrow^S \Gamma_3 \\ (\eta_1 , \eta_2 \mid \sigma_1 \stackrel{S}{\circlearrowleft} \sigma_2) _ x = & \eta_2 \mid (\sigma_1 _ x) [\sigma_2]^S \end{aligned}$$

4.1 Transfer Heaven: Composition of Substitutions with Rewriting

$$\begin{aligned} \text{Compositionality}^{SS} : & \forall (e : \text{Expr } \Gamma_1 T) (\eta_1 : \text{Sub } n_1 n_2) (\eta_2 : \text{Sub } n_2 n_3) \\ & \rightarrow (\sigma_1 : \eta_1 \mid \Gamma_1 \Rightarrow^S \Gamma_2) (\sigma_2 : \eta_2 \mid \Gamma_2 \Rightarrow^S \Gamma_3) \\ & \rightarrow \eta_2 \mid (\eta_1 \mid e [\sigma_1]^S) [\sigma_2]^S \equiv (\eta_1 \circ^S \eta_2) \mid e [\eta_1, \eta_2 \mid \sigma_1 \circ^S \sigma_2]^S \end{aligned}$$
$$\begin{aligned}
& \text{Compositionality}^{SS} (\lambda x. \quad \eta_1 \eta_2 \sigma_1 \sigma_2 = \text{refl}) \\
& \text{Compositionality}^{SS} (\lambda x \{T_1 = T_1\} e) \eta_1 \eta_2 \sigma_1 \sigma_2 = \text{cong } \lambda x (\text{begin} \\
& \quad \eta_2 \mid \eta_1 \mid e [\eta_1 \mid \sigma_1 \uparrow^S T_1]^S [\eta_2 \mid \sigma_2 \uparrow^S (T_1 [\eta_1]^S)]^S \\
& \equiv \langle \text{Compositionality}^{SS} e \eta_1 \eta_2 (\eta_1 \mid \sigma_1 \uparrow^S T_1) (\eta_2 \mid \sigma_2 \uparrow^S (T_1 [\eta_1]^S)) \rangle \\
& \quad (\eta_1 \circ^S \eta_2) \mid e [\eta_1, \eta_2 \mid \eta_1 \mid \sigma_1 \uparrow^S T_1 \circ^S (\eta_2 \mid \sigma_2 \uparrow^S (T_1 [\eta_1]^S))]^S \\
& \equiv \langle \text{cong } ((\eta_1 \circ^S \eta_2) \mid e [\]^S) (\text{Lift-Dist-Comp}^{SS} \sigma_1 \sigma_2) \rangle \\
& \quad (\eta_1 \circ^S \eta_2) \mid e [(\eta_1 \circ^S \eta_2) \mid \eta_1, \eta_2 \mid \sigma_1 \circ^S \sigma_2 \uparrow^S T_1]^S \\
& \square)
\end{aligned}$$
$$\begin{array}{l} \text{Compositionality}^{SS} (e_1 \cdot e_2) \ \eta_1 \ \eta_2 \ \sigma_1 \ \sigma_2 = \text{cong}_2 \text{ } _ _ \\ (\text{Compositionality}^{SS} \ e_1 \ \eta_1 \ \eta_2 \ \sigma_1 \ \sigma_2) \\ (\text{Compositionality}^{SS} \ e_2 \ \eta_1 \ \eta_2 \ \sigma_1 \ \sigma_2) \\ \text{Compositionality}^{SS} (e \cdot^* T) \ \eta_1 \ \eta_2 \ \sigma_1 \ \sigma_2 = \text{cong} \ (_ \cdot^* (T \ [\ \eta_1 \ \circ^S \ \eta_2]^S)) \\ (\text{Compositionality}^{SS} \ e \ \eta_1 \ \eta_2 \ \sigma_1 \ \sigma_2) \end{array}$$

This proof relies on further properties, in particular **Lift-Dist-Comp^{SS}** (lifting a composition of substitutions is equal to composing their liftings) and **Lift*-Dist-Comp^{SS}** (its counterpart for lifting over type bindings), where the statement and proof would require explicit transfers. Essentially, these properties form a hierarchy that stops with lemmas for composing two renamings and composing a renaming with a lifting, all of which work similarly. The supplemental material contains full details.

4.2 Transfer Hell: Composition of Substitutions Without Rewriting

We now contrast the development from Section 4.1 with a proof of the same compositionality property, but carried out with explicit transfers. The proof is taken from the supplement of a paper on a logical relation for System F [10]. We aligned the surface syntax of their definitions with ours to enable direct comparison. Throughout this subsection, code reproduced from that prior work is rendered with a shaded background to distinguish it from our own development.

The type of compositionality already exposes the core issue. Because the σ -calculus equations are only propositional, the theorem statement must include **subst** to bridge indices that are equal but not definitionally equal: the type of the left-hand side of the equality is *not* definitionally equal to the type of the right-hand side. To adapt the types, compositionality of type substitutions is applied to the type of the expression using function **sub**.

Compositionality^{SS} $\equiv$:

$$\begin{aligned} & \forall \{ \eta_1 : \text{Sub } n_1 \ n_2 \} \{ \eta_2 : \text{Sub } n_2 \ n_3 \} \\ & \quad \{ \Gamma_1 : \text{Ctx } n_1 \} \{ \Gamma_2 : \text{Ctx } n_2 \} \{ \Gamma_3 : \text{Ctx } n_3 \} \\ & \quad (e : \text{Expr } n_1 \ \Gamma_1 \ T) \\ & \quad (\sigma_1 : \eta_1 \mid \Gamma_1 \Rightarrow^S \Gamma_2) (\sigma_2 : \eta_2 \mid \Gamma_2 \Rightarrow^S \Gamma_3) \rightarrow \\ & \text{let } \text{sub} = \text{subst } (\text{Expr } n_3 \ \Gamma_3) (\text{compositionality}^{SS} \ T \ \eta_1 \ \eta_2) \text{ in} \\ & \text{sub } (\eta_2 \mid (\eta_1 \mid e \ [\sigma_1]^S) \ [\sigma_2]^S) \equiv (\eta_1 \circ^S \eta_2) \mid e \ [\eta_1, \eta_2 \mid \sigma_1 \circ^S \sigma_2]^S \end{aligned}$$

Because **subst** is part of the statement, every recursive call to the induction hypothesis introduces a fresh application of **subst**, which has to be discharged in the proof alongside the interesting proof steps. To do so, it must be distributed to the outside of the term and combined with the **substs** arising from neighboring subterms before the goal can be closed. This process is further complicated by other occurrences of type substitution lemmas that interact non-trivially with the transports. We illustrate the resulting complexity by zooming in on a single case of the inductive proof—type application—of their compositionality lemma.

Firstly, our development already saves one **subst** in the very definition of substitution application. The type application case requires the swap lemma of type $(T \ [\ T \]^*) \ [\ \eta \] \equiv (T \ [\ \eta \ \uparrow^S \]) \ [\ T \ [\ \eta \] \]^*$ to coerce the term to the expected type, leading to the following definition.

$$\begin{aligned} \eta \mid (_ \cdot _ \{ T = T \} \ e \ T) \ [\ \sigma \]^S &= \text{subst } (\text{Expr } _ _) \\ & \quad (\text{sym } (\text{swap}^{SS} \ \eta \ T \ T)) \\ & \quad ((\eta \mid e \ [\ \sigma \]^S) \cdot^* (T \ [\ \eta \]^S)) \end{aligned}$$

Next, we examine the type application case of the compositionality proof itself. The authors employ heterogeneous equality [8] to avoid some of the administrative reasoning otherwise forced by transfer hell. Without heterogeneous equality, the proof would be even more complex, requiring several auxiliary lemmas that explicitly shuffle applications of **subst** around. Thus, the version below already constitutes a streamlined form of the argument.

```

589 let
590   F2 = Expr n2 Γ2 ; E2 = sym (swapSS η1 T T') ; sub2 = subst F2 E2
591   F3 = Expr n3 Γ3 ; E3 = sym (swapSS (η1 S η2) T T') ; sub3 = subst F3 E3
592   F5 = Expr n3 Γ3 ; E5 = sym (swapSS η2 (T [ η1 S ]) (T' [ η1 ]S)) ; sub5 = subst F5 E5
593 in
594 R.begin
595   η2 | (η1 | (e .* T') [ σ1 ]S) [ σ2 ]S
596 R.≡⟨ refl ⟩ − (1)
597   η2 | (sub2 ((η1 | e [ σ1 ]S) .* (T' [ η1 ]S))) [ σ2 ]S
598 R.≡⟨ H.cong2 {B = Expr n2 Γ2} (λ _ ■ → η2 | ■ [ σ2 ]S) ⟩ − (2)
599   (H.≡-to-≡ (sym E2))
600   (H.≡-subst-removable F2 E2 _) )
601   η2 | ((η1 | e [ σ1 ]S) .* (T' [ η1 ]S)) [ σ2 ]S
602 R.≡⟨ refl ⟩ − (3)
603   sub5 ((η2 | (η1 | e [ σ1 ]S) [ σ2 ]S) .* (T' [ η1 ]S [ η2 ]S))
604 R.≡⟨ H.≡-subst-removable F5 E5 _ ⟩ − (4)
605   (η2 | (η1 | e [ σ1 ]S) [ σ2 ]S) .* (T' [ η1 ]S [ η2 ]S)
606 R.≡⟨ H.cong3 {B = λ ■ → Expr n3 Γ3 (∀α ■)} {C = λ _ _ → Type n3 _} − (5)
607   (λ _ ■1 ■2 → ■1 .* ■2)
608   (H.≡-to-≡ (lift-compositionalitySS T η1 η2))
609   (CompositionalitySS e σ1 σ2)
610   (H.≡-to-≡ (compositionalitySS T' η1 η2)) )
611   ((η1 S η2) | e [ η1 , η2 | σ1 S σ2 ]S) .* (T' [ η1 S η2 ]S)
612 R.≡⟨ H.sym (H.≡-subst-removable F3 E3 _) ⟩ − (6)
613   sub3 (((η1 S η2) | e [ η1 , η2 | σ1 S σ2 ]S) .* (T' [ η1 S η2 ]S))
614 R.≡⟨ refl ⟩ − (7)
615   (η1 S η2) | (e .* T') [ η1 , η2 | σ1 S σ2 ]S
616 R.□

```

The proof proceeds in seven steps.

- (1) Unfold the definition of expression substitution at the inner application of σ_1 , exposing the **subst** inside.
- (2) Remove the **subst** that sits inside the term under consideration, justified by heterogeneous equality.
- (3) Reveal yet another **subst** introduced by the substitution application of σ_2
- (4) Remove the **subst** on the outside.
- (5) Apply the induction hypothesis while simultaneously coercing the index according to the compositionality law for type substitutions.
- (6) Reinsert the **subst** demanded by the proof goal on the outside.
- (7) Fold the definition of expression substitution to conclude the proof.

Of these, steps (2), (4), the type coercion in (5), and (6), as well as the use of heterogeneous equality, are entirely absent from our rewriting-based proof, only the application of the inductive hypothesis and the implicit (un)folder of definitions remain.

	With Rewriting		Transfer Hell	
	LOC	Chars	LOC	Chars
Type substitution + proofs	217	12 246	285	12 584
Expression substitution + proofs	199	11 465	1 405	69 846
Total	416	23 711	1 690	82 430
Type-check time	4.49 s		9.51 s	

Table 2. Proof statistics

4.3 Quantitative Comparison

We complement the preceding qualitative comparison with a quantitative one. Table 2 records the lines of code and number of characters required to mechanize the compositionality law for expression substitutions, together with the type-checking time of the respective developments.

As expected, the amount of code required to formalize type substitution is roughly the same in both developments. For expression substitutions, however, the picture changes: the transfer-hell version requires more than five times as many lines of code as the rewriting-based one. As a welcome side effect, the type-checking time is also reduced by about 50% when the σ -calculus equations are installed as rewrite rules.

5 Methodology

5.1 Safety Conditions for Rewriting

Rewrite rules and the reductions of a dependently typed system interact in non-trivial ways [3, 6, 7]. According to RTT's criteria, we can make the equations from Section 2.3 definitional because they preserve subject reduction and consistency.

To uphold subject reduction, the critical property of the newly added rewriting equations is that their *union with the existing reduction relation* is confluent. Agda already provides a standard set of confluent reduction rules: β , η (function and record expansion), δ (definition unfolding), and ι (dependent pattern matching clauses). Agda checks confluence of the extended system using a conservative set of simple syntactic criteria [4].

Some β and ι reductions arising from the functional definitions of renamings and substitutions in Figure 3 interfere with the intended rewrite equalities so that confluence breaks. For this reason, we explicitly hide these definitions inside an `opaque` block, so that only the equations we install with `REWRITE` define their reduction behaviour.

Without the `opaque` definitions, it may happen that the left-hand side of an equation reduces to a normal form, which is different from the right-hand side. Thus, confluence would be violated. As a concrete example, consider the equation η -law (left) in connection with the definition of composition (right):

$$(\text{zero} \&^S \eta) \cdot^S (\langle \text{wk}^R \rangle \circ^S \eta) \equiv \eta \qquad (\eta_1 \circ^S \eta_2) \alpha = (\eta_1 \alpha) [\eta_2]^S$$

Without the `opaque` block, the definition of composition would δ -unfold the inner application $\langle \text{wk}^R \rangle \circ^S \eta$ to $\lambda \alpha \rightarrow \eta (\text{succ } \alpha)$ so that the left-hand side reduces as follows:

$$\begin{aligned} (\text{zero} \&^S \eta) \cdot^S (\langle \text{wk}^R \rangle \circ^S \eta) &\equiv (\text{zero} \&^S \eta) \cdot^S (\lambda \alpha \rightarrow \eta (\text{succ } \alpha)) \\ &\equiv (\eta \text{ zero}) \cdot^S (\lambda \alpha \rightarrow \eta (\text{succ } \alpha)) \end{aligned}$$

Two issues arise. First, the definitional reduction happens *before* the rewrite engine tries to match the left-hand side of η -law syntactically. Thus the rewrite never happens³. Second, if the rewrite rule were admitted, confluence would fail. The two reducts η (via the rewrite) and $(\eta \text{ zero}) \cdot^S (\lambda \alpha \rightarrow \eta (\text{succ } \alpha))$ (via δ -unfolding) are propositionally equal but not definitionally equal. Because the function $\cdot^S$ case-splits on a neutral variable, no further reduction back to η is possible. Hiding the definition of the symbol η^S (together with $\cdot^S$ and the other operators) inside an **opaque** block turns them into stuck symbols, restoring both matchability and confluence.

Consistency is preserved as long as the equalities we rewrite are instantiated with proofs. This is easy to see, because RTT can be translated to extensional type theory [4], a theory known to be consistent.

Agda currently supports all features we require: opaque-defined symbols, proof-instantiated rewrite equations, and triangle-based confluence checking. Rocq currently lacks the first two, although the overall method otherwise transfers.

After automating the type-level substitution theory by installing safe rewrites, we aim to extend our method to the expressions level.

5.2 Equational Theory for Expression Substitution

With functoriality for expression substitution in place, one might ask whether we can go further and install a full equational theory for expression substitution as definitional equalities, just as we did at the type level. Indeed, the σ -calculus for type-substitution-indexed expression substitutions is an extension of the one for type substitutions in Section 2.3, and the same holds analogously for expression renamings.

Expression substitutions carry additional operators that have no counterpart at the type level. Every law involving extension or weakening splits into two expression-level laws. One governs the expression variants $(\cdot^S, \text{wk}^S)$ and directly mirrors the type-level rule Section 2.3. The other governs the type variants $(\cdot^{S*}, \text{wk}^{S*})$ and is genuinely new. We illustrate this pattern with the η -id law. The expression variant

$$\eta\text{-Id}^S : \langle \text{id}^R \rangle \mid (\text{zero}) \cdot^S (\text{Wk}^S T) \equiv \text{Id}^S \{ \Gamma = \Gamma \triangleright T \}$$

is a direct transcription of η -id from Section 2.3, whereas the type variant

$$\eta\text{-Id}^{S*} : \langle \text{wk}^R \rangle \mid (\text{zero}) \cdot^{S*} \text{wk}^{S*} \equiv \text{Id}^S \{ \Gamma = \Gamma \triangleright * \}$$

asserts that extending the type-binding weakening with the topmost type variable zero yields the identity expression substitution. The same doubling applies to the remaining laws in which extension or weakening appears (η , interact, distributivity and some β laws).

The full equational system is included in the supplement, and Agda verifies its confluence. For the confluence check to go through, we had to artificially block reduction of the traversal functions $\llbracket _ \rrbracket^R$ and $\llbracket _ \rrbracket^S$ and re-install each of their defining clauses as a rewrite rule. In the type application case, rewriting must perform two reduction steps at once to satisfy the triangle criterion. One reduction normalizes the type index; the other pushes substitution into the expression.

5.3 Practical Guidelines for Other Intrinsic Developments

We have now made the equations of the σ -calculus definitional at two levels of the syntactic hierarchy: at the type level and at the expression level, where the latter depends on the former. In general, this strategy applies to arbitrarily tall towers of intrinsically typed syntaxes, in which each level's type indices themselves support substitution.

³To see why, note that η does not appear in pattern position in the reduct. A precise definition of pattern position is available in the Agda Wiki: <https://agda.readthedocs.io/en/latest/language/rewriting.html#general-shape-of-rewrite-rules>

More precisely, consider a tower $T_0, T_1, \dots, T_n$ of syntactic sorts in which the variables and substitutions of T_{i+1} are indexed by the contexts and substitutions of $T_0, \dots, T_i$. The recipe proceeds bottom-up:

Methodology Recipe. To avoid transfer hell in intrinsically typed towers:

- (1) Start at the base sort T_0 : define renamings/substitutions and install the σ -calculus equations as rewrite rules.
- (2) For each higher sort T_i , define syntax and substitution *on top of* already-definitional lower-level substitutions.
- (3) Prove the σ -calculus laws for T_i and install them as rewrite rules.
- (4) Check confluence of the combined reduction relation after each level.

The key is that the laws for all levels $< i$ are already definitional, when we define and prove the laws for level i . Otherwise, the types of the level- i operations would no longer match definitionally, and we would be back in transfer hell.

The size of the equational theory grows with each level. A substitution at level i can be lifted under a binder of any sort $T_0, \dots, T_{i-1}$, and each kind of binder induces its own family of laws, for weakening and extension, and the corresponding interactions with composition and traversal. At the type level this contributes no additional laws. At the expression level we already saw a few new laws, namely the laws governing weakening and extension by types. At level three one would add even more laws, and so on. The rule count therefore grows linearly with the level, but each new family follows the same template as the previous one, so the proof effort per family stays constant.

As an example from the literature, the language System $F_C^\dagger$ [14] features three levels: kinds, types, and expressions. Each level carries its own variables and substitution. Using our approach, it should be possible to construct an intrinsically typed representation for it and establish the substitution theory for all three levels uniformly, without any transfer reasoning.

6 Conclusion

Intrinsically typed syntax facilitates elegant mechanizations by making well-typed programs representable only as well-typed data. For polymorphic calculi this promise is often undermined by pervasive transport reasoning. Our experience with System F shows that the difficulty does not stem from binding structure or indexing itself, but from a mismatch between substitution and definitional equality: substitutions that are extensionally equal fail to compute to a canonical form. By internalizing the equational laws of the σ -calculus as definitional equalities, substitutions become canonical and routine transports disappear from proofs. The resulting developments recover the elegance and concision seen in the simply-typed case while scaling to towers of syntactic sorts. In our System F case study, we observe a decrease of code size by a factor of 3–4 and a speedup by a factor of 2. More broadly, this fresh perspective proposes a practical recipe for mechanizing type systems in dependently typed proof assistants: when intrinsically typed encodings become proof-heavy, a promising strategy is often to strengthen definitional equality rather than introduce additional lemmas or automation.

References

- [1] Nick Benton, Chung-Kil Hur, Andrew Kennedy, and Conor McBride. 2012. Strongly Typed Term Representations in Coq. *J. Autom. Reason.* 49, 2 (2012), 141–159. doi:10.1007/S10817-011-9219-0
- [2] James Chapman, Roman Kireev, Chad Nester, and Philip Wadler. 2019. System F in Agda, for Fun and Profit. In *Mathematics of Program Construction: 13th International Conference, MPC 2019, Porto, Portugal, October 7–9, 2019, Proceedings* (Porto, Portugal). Springer-Verlag, Berlin, Heidelberg, 255–297. doi:10.1007/978-3-030-33636-3_10
- [3] Jesper Cockx. 2020. Type Theory Unchained: Extending Agda with User-Defined Rewrite Rules. In *25th International Conference on Types for Proofs and Programs (TYPES 2019) (Leibniz International Proceedings in Informatics (LIPIcs))*.

- Vol. 175), Marc Bezem and Assia Mahboubi (Eds.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 2:1–2:27. [doi:10.4230/LIPICs.TYPES.2019.2](https://doi.org/10.4230/LIPICs.TYPES.2019.2)
- [4] Jesper Cockx, Nicolas Tabareau, and Théo Winterhalter. 2021. The Taming of the Rew: A Type Theory with Computational Assumptions. *Proc. ACM Program. Lang.* 5, POPL, Article 60 (Jan. 2021), 29 pages. [doi:10.1145/3434341](https://doi.org/10.1145/3434341)
- [5] Wen Kokke, Jeremy G. Siek, and Philip Wadler. 2020. Programming Language Foundations in Agda. *Sci. Comput. Program.* 194 (2020), 102440. [doi:10.1016/J.SCICO.2020.102440](https://doi.org/10.1016/J.SCICO.2020.102440)
- [6] Yann Leray, Gaëtan Gilbert, Nicolas Tabareau, and Théo Winterhalter. 2024. The Rewster: Type Preserving Rewrite Rules for the Coq Proof Assistant. In *15th International Conference on Interactive Theorem Proving (ITP 2024) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 309)*, Yves Bertot, Temur Kutsia, and Michael Norrish (Eds.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 26:1–26:18. [doi:10.4230/LIPICs.ITP.2024.26](https://doi.org/10.4230/LIPICs.ITP.2024.26)
- [7] Yann Leray and Théo Winterhalter. 2026. Encode the Cake and Eat It Too: Controlling Computation in Type Theory, Locally. *Proc. ACM Program. Lang.* 10, POPL, Article 62 (Jan. 2026), 30 pages. [doi:10.1145/3776704](https://doi.org/10.1145/3776704)
- [8] Conor McBride. 2000. Elimination with a Motive. In *Types for Proofs and Programs, International Workshop, TYPES 2000, Durham, UK, December 8-12, 2000, Selected Papers (Lecture Notes in Computer Science, Vol. 2277)*, Paul Callaghan, Zhaohui Luo, James McKinna, and Robert Pollack (Eds.). Springer, 197–216. [doi:10.1007/3-540-45842-5_13](https://doi.org/10.1007/3-540-45842-5_13)
- [9] Conor McBride. 2005. Type-preserving Renaming and Substitution. (2005). <http://strictlypositive.org/ren-sub.pdf> Unpublished manuscript.
- [10] Hannes Saffrich, Peter Thiemann, and Marius Weidner. 2024. Intrinsically Typed Syntax, a Logical Relation, and the Scourge of the Transfer Lemma. In *Proceedings of the 9th ACM SIGPLAN International Workshop on Type-Driven Development (Milan, Italy) (TyDe 2024)*. Association for Computing Machinery, New York, NY, USA, 2–15. [doi:10.1145/3678000.3678201](https://doi.org/10.1145/3678000.3678201)
- [11] Steven Schäfer, Tobias Tebbi, and Gert Smolka. 2015. Autosubst: Reasoning with de Bruijn Terms and Parallel Substitutions. In *Interactive Theorem Proving*, Christian Urban and Xingyuan Zhang (Eds.). Springer International Publishing, Cham, 359–374. https://www.ps.uni-saarland.de/Publications/documents/SchaeferEtAl_2015_Autosubst_Reasoning.pdf
- [12] Kathrin Stark. 2020. *Mechanising Syntax with Binders in Coq*. Ph.D. Dissertation. Saarland University, Saarbrücken, Germany. https://www.ps.uni-saarland.de/Publications/documents/Stark_2020_Mechanising.pdf
- [13] Kathrin Stark, Steven Schäfer, and Jonas Kaiser. 2019. Autosubst 2: Reasoning with Multi-sorted de Bruijn Terms and Vector Substitutions (*CPP 2019*). Association for Computing Machinery, New York, NY, USA, 166–180. [doi:10.1145/3293880.3294101](https://doi.org/10.1145/3293880.3294101)
- [14] Brent A. Yorgey, Stephanie Weirich, Julien Cretin, Simon Peyton Jones, Dimitrios Vytiniotis, and José Pedro Magalhães. 2012. Giving Haskell a Promotion. In *Proceedings of the 8th ACM SIGPLAN Workshop on Types in Language Design and Implementation (Philadelphia, Pennsylvania, USA) (TLDI '12)*. Association for Computing Machinery, New York, NY, USA, 53–66. [doi:10.1145/2103786.2103795](https://doi.org/10.1145/2103786.2103795)